

18: Application Example OCR

[Previous](#) [Next](#) [Index](#)

Problem description and pipeline

- Case study focused around photo OCR
- Three reasons to do this
 - 1) Look at how a **complex system** can be put together
 - 2) The idea of a machine learning **pipeline**
 - What to do next
 - How to do it
 - 3) Some more interesting ideas
 - Applying machine learning to tangible problems
 - **Artificial data synthesis**

What is the photo OCR problem?

- Photo OCR = photo optical character recognition
 - With growth of digital photography, lots of digital pictures
 - One idea which has interested many people is getting computers to understand those photos
 - The photo OCR problem is getting computers to read text in an image
 - Possible applications for this would include
 - Make searching easier (e.g. searching for photos based on words in them)
 - Car navigation
- OCR of documents is a comparatively easy problem
 - From photos it's really hard

OCR pipeline

- 1) Look through image and find text
- 2) Do character segmentation
- 3) Do character classification
- 4) *Optional* some may do spell check after this too
 - We're not focussing on such systems though



- **Pipelines** are common in machine learning
 - Separate modules which may each be a machine learning component or data processing component
- If you're designing a machine learning system, pipeline design is one of the most important questions
 - Performance of pipeline and each module often has a big impact on the overall performance a problem
 - You would often have different engineers working on each module
 - Offers a natural way to divide up the workload

Sliding window image analysis


- How do the individual models work?
- Here focus on a sliding windows classifier
- As mentioned, stage 1 is **text detection**
 - Unusual problem in computer vision - different rectangles (which surround text) may have different aspect ratios (aspect ratio being height : width)
 - Text may be short (few words) or long (many words)
 - Tall or short font
 - Text might be straight on
 - Slanted



- Let's start with a simpler example

Pedestrian detection

- Want to take an image and find pedestrians in the image



- This is a slightly simpler problem because the aspect ration remains pretty constant
- Building our detection system
 - Have 82 x 36 aspect ratio
 - This is a typical aspect ratio for a standing human
 - Collect training set of positive and negative examples




Positive examples ($y = 1$) Negative examples ($y = 0$)

- Could have 1000 - 10 000 training examples
- Train a neural network to take an image and classify that image as pedestrian or not
 - Gives you a way to train your system
- Now we have a new image - how do we find pedestrians in it?
 - Start by taking a rectangular 82 x 36 patch in the image



- Run patch through classifier - hopefully in this example it will return $y = 0$
- Next slide the rectangle over to the right a little bit and re-run
 - Then slide again
 - The amount you slide each rectangle over is a parameter called the step-size or stride
 - Could use 1 pixel
 - Best, but computationally expensive
 - More commonly 5-8 pixels used
 - So, keep stepping rectangle along all the way to the right
 - Eventually get to the end
 - Then move back to the left hand side but step down a bit too
 - Repeat until you've covered the whole image
- Now, we initially started with quite a small rectangle
 - So now we can take a larger image patch (of the same aspect ratio)
 - Each time we process the image patch, we're resizing the larger patch to a smaller image, then running that smaller image through the classifier
- Hopefully, by changing the patch size and rastering repeatedly across the image, you eventually recognize all the pedestrians in the picture




Text detection example

- Like pedestrian detection, we generate a labeled training set with
 - Positive examples (some kind of text)
 - Negative examples (not text)



Positive examples ($y = 1$) Negative examples ($y = 0$)


- Having trained the classifier we apply it to an image
 - So, run a sliding window classifier at a fixed rectangle size
 - If you do that end up with something like this



- White region show where text detection system thinks text is
 - Different shades of gray correspond to probability associated with how sure the classifier is the section contains text
 - Black - no text
 - White - text
 - For text detection, we want to draw rectangles around all the regions where there is text in the image
- Take classifier output and apply an **expansion algorithm**
 - Takes each of white regions and expands it
 - How do we implement this
 - Say, for every pixel, is it within some distance of a white pixel?
 - If yes then colour it white




- Look at connected white regions in the image above
 - Draw rectangles around those which make sense as text (i.e. tall thin boxes don't make sense)




- This example misses a piece of text on the door because the aspect ratio is wrong
 - Very hard to read

Stage two is character segmentation

- Use supervised learning algorithm
- Look in a defined image patch and decide, is there a split between two characters?
 - So, for example, our first training data item below looks like there is such a split
 - Similarly, the negative examples are either empty or hold a full characters



- We train a classifier to try and classify between positive and negative examples
 - Run that classifier on the regions detected as containing text in the previous section
- Use a 1-dimensional sliding window to move along text regions
 - Does each window snapshot look like the split between two characters?
 - If yes insert a split
 - If not move on
 - So we have something that looks like this



Character classification

- Standard OCR, where you apply standard supervised learning which takes an input and identify which character we decide it is
 - Multi-class characterization problem

Getting lots of data: Artificial data synthesis

- We've seen over and over that one of the most reliable ways to get a high performance machine learning system is to take a low bias algorithm and train on a massive data set
 - Where do we get so much data from?
 - In ML we can do artificial data synthesis
 - This doesn't apply to every problem
 - If it applies to your problem, it can be a great way to generate loads of data
- Two main principles
 - 1) Creating data from scratch
 - 2) If we already have a small labeled training set can we amplify it into a larger training set

Character recognition as an example of data synthesis

- If we go and collect a large labeled data set will look like this
 - The goal is to take an image patch and have the system recognize the character
 - Let's treat the images as gray-scale (makes it a bit easier)




Real data

- How can we amplify this
 - Modern computers often have a big font library
 - If you go to websites, huge free font libraries
 - For more training data, take characters from different fonts, paste these characters again random backgrounds
- After some work, can build a synthetic training set



Synthetic data

- Random background
- Maybe some blurring/distortion filters
- Takes thought and work to make it look realistic
 - If you do a sloppy job this won't help!
 - So unlimited supply of training examples
- This is an example of creating new data from scratch
- Other way is to introduce distortion into existing data
 - e.g. take a character and warp it



- 16 new examples
- Allows you amplify existing training set
- This, again, takes thought and insight in terms of deciding how to amplify

Another example: speech recognition

- Learn from audio clip - what were the words
 - Have a labeled training example
 - Introduce audio distortions into the examples
- So only took one example
 - Created lots of new ones!
- When introducing distortion, they should be reasonable relative to the issues your classifier may encounter

Getting more data

- Before creating new data, make sure you have a low bias classifier
 - Plot learning curve
- If not a low bias classifier increase number of features
 - Then create large artificial training set
- Very important question: How much work would it be to get 10x data as we currently have?
 - Often the answer is, "Not that hard"
 - This is often a huge way to improve an algorithm
 - Good question to ask yourself or ask the team
- How many minutes/hours does it take to get a certain number of examples
 - Say we have 1000 examples
 - 10 seconds to label an example
 - So we need another 9000 - 90000 seconds
 - Comes to a few days (25 hours!)
- Crowd sourcing is also a good way to get data
 - Risk or reliability issues
 - Cost
 - Example
 - E.g. Amazon mechanical turks

Ceiling analysis: What part of the pipeline to work on next


- Through the course repeatedly said one of the most valuable resources is developer time
 - Pick the right thing for you and your team to work on
 - Avoid spending a lot of time to realize the work was pointless in terms of enhancing performance

Photo OCR pipeline

- Three modules
 - Each one could have a small team on it
 - Where should you allocate resources?
- Good to have a single real number as an evaluation metric
 - So, character accuracy for this example
 - Find that our test set has 72% accuracy

Ceiling analysis on our pipeline


- We go to the first module
 - Mess around with the test set - manually tell the algorithm where the text is
 - Simulate if your text detection system was 100% accurate
 - So we're feeding the character segmentation module with 100% accurate data now
 - How does this change the accuracy of the overall system



- Accuracy goes up to 89%
- Next do the same for the character segmentation
 - Accuracy goes up to 90% now
- Finally do the same for character recognition
 - Goes up to 100%
- Having done this we can qualitatively show what the upside to improving each module would be
 - Perfect text detection improves accuracy by 17%!
 - Would bring the biggest gain if we could improve
 - Perfect character segmentation would improve it by 1%
 - Not worth working on
 - Perfect character recognition would improve it by 10%
 - Might be worth working on, depends if it looks easy or not
- The "ceiling" is that each module has a ceiling by which making it perfect would improve the system overall

Other example - face recognition

- NB this is not how it's done in practice



- Probably more complicated than is used in practice
- How would you do ceiling analysis for this
 - Overall system is 85%
 - + Perfect background -> 85.1%
 - Not a crucial step
 - + Perfect face detection -> 91%
 - Most important module to focus on
 - + Perfect eyes -> 95%
 - + Perfect nose -> 96%
 - + Perfect mouth -> 97%
 - + Perfect logistic regression -> 100%
- Cautionary tale
 - Two engineers spent 18 months improving background pre-processing
 - Turns out had no impact on overall performance
 - Could have saved three years of man power if they'd done ceiling analysis