

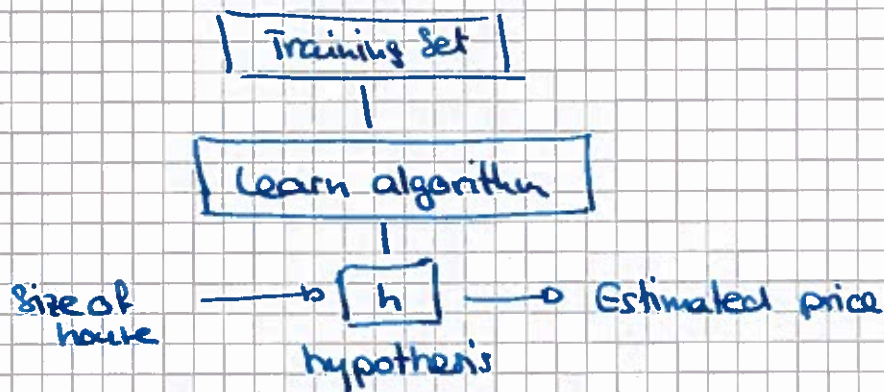
2. LINEAR REGRESSION

Chapter 2

Linear Regression

Notation:

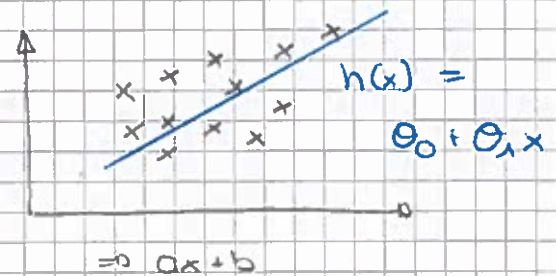
- m - number of training examples
- x - input variable / feature
- y - output variable
- (x, y) simple training example
- $(x^{(i)}, y^{(i)})$ specific example



Representation of h

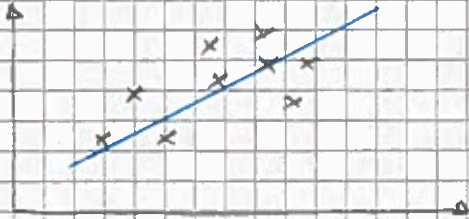
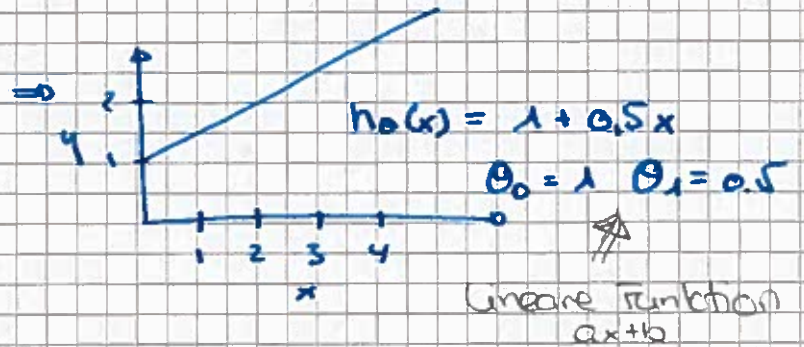
$$\Rightarrow h_{\theta}(x) = \theta_0 + \theta_1 x \Rightarrow$$

↳ shorthand $h(x)$



Linear Regression with one variable
Univariate Linear Regression

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$



Idea: Choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for training example (x, y)

minimize θ_0, θ_1

$$\frac{1}{2m} \sum_{i=1}^m \underbrace{(h_{\theta}(x^{(i)}))}_{\text{Predicted value}} - \underbrace{y^{(i)}}_{\text{real value}})^2$$

of training examples

\Downarrow

$$h(x^{(i)}) = \theta_0 + \theta_1 x \Rightarrow \text{given by the data}$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

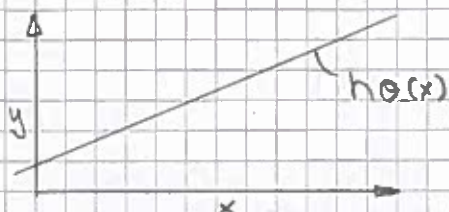
minimize $= J(\theta_0, \theta_1) \Rightarrow$ Cost Function

\Rightarrow Squared error function

Summary:

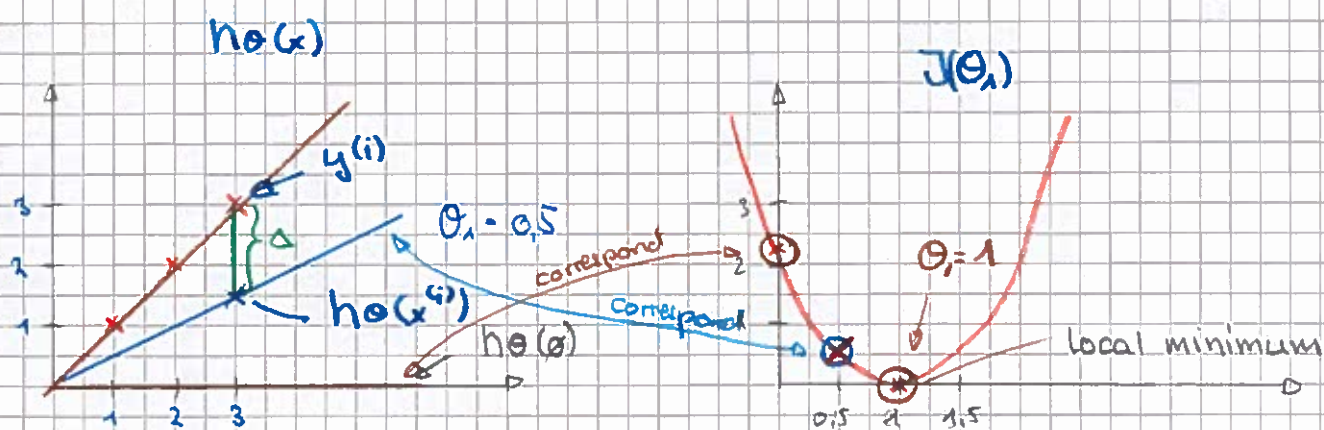
Hypothesis: $h(x) = \theta_0 + \theta_1 x$

Parameters: θ_0, θ_1



Cost function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$

Goal: minimize $J(\theta_0, \theta_1)$



$$J(0.5) = \frac{1}{2m} \sum_{i=1}^m = \frac{1}{2 \cdot 3} [(0.5-1)^2 + (1-2)^2 + (1.5-3)^2]$$

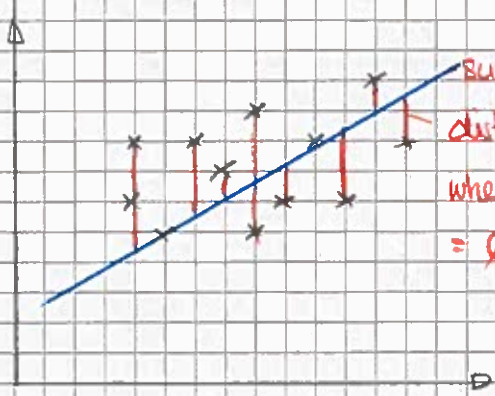
$$= \frac{1}{2 \cdot 3} \cdot 3.5 = 0.58$$

Δ difference between predicted and effective value

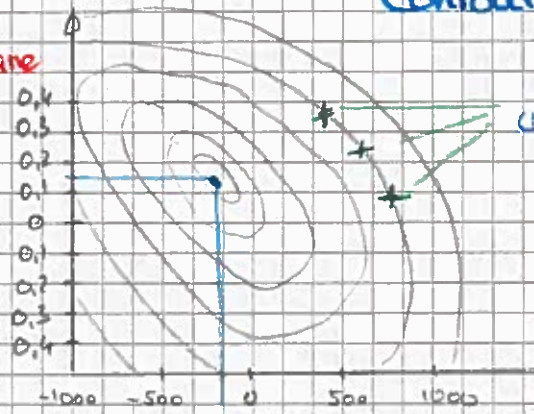
minimize $J(\theta_1)$

$$J(0) = \frac{1}{2m} \cdot [(0-1)^2 + (0-2)^2 + (0-3)^2] = \frac{1 \cdot 14}{2 \cdot 3} = \sim 2.3$$

$$J(1) = \frac{1}{2m} \cdot [(1-1)^2 + (2-2)^2 + (3-3)^2] = \frac{0}{6} = 0$$



Sum of square
distance
where error
= 0



Contour plots

all same distance

minimized value

Gradient descent algorithm

5

=> want to get the local/global minimum

$$\min J(\theta_0, \theta_1)$$

=> Gradient descent applies to more general functions

$$J(\theta_0, \theta_1, \dots, \theta_n) \Rightarrow \min J(\theta_0, \theta_1, \dots, \theta_n)$$

Repeat until convergence ϵ

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j=0, j=1)$$

}

=> simultaneous update

$$\text{temp}0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

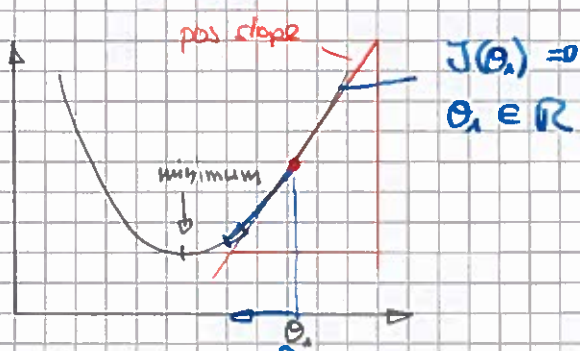
α = learning rate

$$\text{temp}1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp}0$$

$$\theta_1 := \text{temp}1$$

derivative term

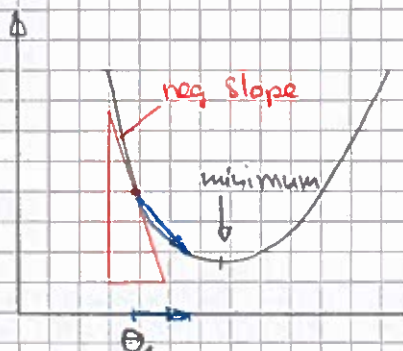


$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

≥ 0

$$\theta_1 := \theta_1 - \alpha \cdot (\text{pos. number})$$

always pos.



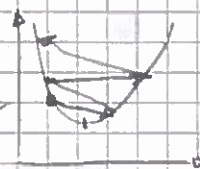
$$\frac{\partial}{\partial \theta_1} J(\theta_1)$$

≤ 0

$$\theta_1 := \theta_1 - \alpha \cdot (\text{neg. number})$$

make sure that α is not too small or too large

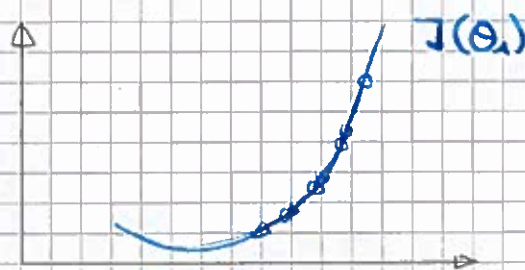
- too small might take too many steps to converge
- too big might overshoot the min. and might never converge.



Gradient descent algorithm

16

Gradient descent can converge to a local minimum, even with the learning rate α fixed.



$$\theta_1 := \theta_1 - \alpha \underbrace{\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)}_x$$

As we approach local minimum, gradient descent will take smaller steps.

term x will decrease therefore value will be smaller & smaller

No need to decrease α .

Linear Regression with One Variable

Apply gradient descent to linear regression function

$$\begin{aligned} \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) &= \frac{\partial}{\partial \theta_1} \cdot \frac{1}{2m} \cdot \sum_{i=1}^m (\underbrace{h_{\theta}(x^{(i)})}_{\theta_0 + \theta_1 x^{(i)}} - y^{(i)})^2 \\ &= \frac{\partial}{\partial \theta_1} \cdot \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 \end{aligned}$$

$$\theta_0 := \theta_0 \quad j=0 : \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 \quad j=1 : \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

Gradient descent algorithm

repeat until convergence $\{$

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

$\}$

always update simultaneously.